I, ME AND MYSELF !!!

Home About Me

MONDAY, SEPTEMBER 14, 2009

Segmented Sieve

Memory and time efficient:)

Problem Statement:

#include <string.h>

Your are given two integers a and b. You have to find all the primes within range a and b. Here, $1 \le a \le b \le 2^{31}$ -1 and b - $a \le 10^5$.

Note: You have to handle 1, 2 and even numbers for appropriate case of your own.

Solution:

```
#define MAX 46656
#define LMT 216
#define LEN 4830
#define RNG 100032
unsigned base[MAX/64], segment[RNG/64], primes[LEN];
#define sq(x) ((x)*(x))
#define mset(x,v) memset(x,v,sizeof(x))
#define chkC(x,n) (x[n>>6]&(1<<((n>>1)&31)))
#define setC(x,n) (x[n>>6]|=(1<<((n>>1)&31)))
/* Generates all the necessary prime numbers and marks them in base[]*/
void sieve()
    unsigned i, j, k;
    for(i=3; i<LMT; i+=2)</pre>
        if(!chkC(base, i))
            for(j=i*i, k=i<<1; j<MAX; j+=k)
                setC(base, j);
    for(i=3, j=0; i<MAX; i+=2)
        if(!chkC(base, i))
            primes[j++] = i;
}
/* Returns the prime-count within range [a,b] and marks them in segment[]
*/
int segmented_sieve(int a, int b)
    unsigned i, j, k, cnt = (a \le 2 \&\& 2 \le b)? 1 : 0;
    if(b<2) return 0;
    if(a<3) a = 3;
    if(a%2==0) a++;
    mset(segment,0);
    for(i=0; sq(primes[i]) \le b; i++)
```



Zobayer Hasan Dhaka, Bangladesh Undergraduate Student School: CSE-DU

View my complete profile

SUBSCRIBE

▶ Posts▶ Comments

BLOG HITS



BLOG ARCHIVE

- **▶** 2011 (15)
- **2010** (33)
- ▼ 2009 (27)
 - ► December (11)
 - ► November (4)
 - October (1)
- ▼ September (2)

 SPOJ Solve list comparison tool

 Segmented Sieve
- ► August (1)
- ▶ July (8)

CATAGORIES

academic study (15) access modifiers (1) algorithm (28) bash (1) beginner (17) bfs (1) bigint (1) binary tree (1) bitwise (4) blogger (5) bpm (2) brainfuck (1) bst (1) c (1) C++ (36) changes (1) character device driver (1) combinatorics (2) comparator (1) compression (1) computational geometry (2)

```
{
    j = primes[i] * ( (a+primes[i]-1) / primes[i] );
    if(j%2==0) j += primes[i];
    for(k=primes[i]<<1; j<=b; j+=k)
        if(j!=primes[i])
        setC(segment, (j-a));
}
for(i=0; i<=b-a; i+=2)
    if(!chkC(segment, i))
        cnt++;
return cnt;
}</pre>
```

This is a sample program which demonstrates segmented sieve. Very fast and memory efficient version. 'base' is the array which holds the flags for all the primes upto $\sqrt{(2^{31}\text{-}1)}$, i.e. the square-root of the max limit, and all the primes are stored in the 'primes' array. Later, these primes are used to determine whether a number is a composite or not within a certain range in the segmented sieve. To avoid overflow and sign bit problems, unsigned type is used.

A little explanation:

First of what what these macros mean?
#define MAX 46656
#define LMT 216
#define LEN 4830
#define RNG 100032

MAX is the sqrt of maximum possible input, in case of here, the maximum is integer range sqrt of which is almost MAX used here. So, MAX is not maximum allowed input, it is just sqrt of maximum input which is pretty big as 2147483647 i.e. 32 bit signed integer maximum.

LMT is sqrt of MAX. We all know, we run sieve upto sqrt MAX

LEN is the maximum possible different primes that can be stored using this algorithm with specific range defined as **RNG**, on which the segmented sieve will run and collect the primes out of it.

Now the next two vital macros:
#define chkC(x,n) (x[n>>6]&(1<<((n>>1)&31)))
#define setC(x,n) (x[n>>6]|=(1<<((n>>1)&31)))
And why we divide by 64:

Ok, yes, it is clearly bit shifting. But you must know what we do in bitwise sieve first in order to get this. Instead of using a whole array position to store just one flag, we can use its each 32 bits to store one flag, which saves memory by a factor 1/32. So its very logical to capture memory upto MAX/32, but why MAX/64 and RNG/64 here?

Because, we really have no point of handling the even number as 2 is the only even prime and we can handle it manually, without any stress. So, if we do not consider the even numbers at all, the total numbers are again reduced by a factor 1/2, isn't it? So what we get total is MAX/32/2 = MAX/64, same for RNG.

Now, the two macros chkC and setC is pretty straight forward. chkC checks if a specific bitflag is 1 or 0, and setC sets a specific bitflag 1 to mark it as a composite. They work similarly, so I will explain only the chkC part.

In bitwise sieve, where is a specific value n located? Obviously $(n/32)^{th}$ index, and on that index, $(n\%32)^{th}$ bit from LSB (right hand side). But we just said before, we are interested with only odd numbers, so we map n with n/2 as follows:

confusion (1) contest (4) crc (1) cse (3) css (1) customize (1) data structure (10) database (1) decoding (1) design (1) device driver (1) divide and conquer (2) dp (3) driver (1) dynamic programming (8) encoding (1) encryption (1) error (2) esoteric language (2) euler circuit (1) euler path (1) expression evaluation (1) extended euclid (1) facebook (1) factorization (1) funny (14) gcd (2) geometry (4) graph (7) hashing (1) hints (5) hopcroft karp (1) huffman (1) java (4) javascript (1) jdbc (1) kernel programming (2) lab (1) like (1) linear algebra (3) linux (5) ls (1) makefile (1) math (16) matrix (2) matrix algebra (1) matrix exponentiation (1) matrix multiplication (1) maxflow (1) maximum bipartite matching (2) maximum flow (1) merge sort (3) mistake (1) modular arithmatic (1) module compiling (2) mysql (1) number system (1) number theory (8) online judge (3) operating system (1) os (1) other (8) parallel programming (1) pollard rho (1) primes (3) problem (3) problem classifier (2) problem solving (31) programming (48) pthreaded (1) puzzle (1) python (3) recursion (5) shell (1) shell script (1) sieve (1) simulation (1) sort (3) spacing (1) sphere online judge (12) spoj (12) syntax highlighting (1) system programming (4) table tag (1) tc (1) template (4) thread (1) topcoder (1) tree (1) ubuntu (1) usaco (1) uva (5) uva online judge

(5) vector (1)

So, n is actually n/2, which implies the previous statement: n's (actually n/2 's) position is in index [n/2/32] = [n/64] = [n>>6] and on the bit position (n/2)%32 = (n>>1)&31; [We know, modding with a power of 2 is same as ANDing with (same power of 2)-1]. The rest is, how we check / set this specific bit. I have another post explaining these operations: Bitwise operations in C: Part 3, and obviously you can google it:)



<3 th!2 1st.

Posted by Zobayer Hasan at 1:59 AM

10 comments:

e fushar said...

Thanks! After reading this article I managed to get AC on SPOJ's PRIME1.

September 27, 2009 10:42 AM

Arun Chauhan said...

great code dude...

but i have some problems understanding it....

if you don't mind can you explain me.... some of the following things....

i want help in understanding the

#define chkC(x,n) (x[n>>6]&(1<<((n>>1)&31)))

#define setC(x,n) (x[n>>6]|=(1<<((n>>1)&31)))

i know they are using bitshift and operators..... but what exactly they are trying to do can

somebody make me understand.... why is it needed to divide n by 64 and why are they ANDing (n/2) with 31..... and why MAX=46656 and RNG=100032 and LEN=4830....

LMT is ok b'cos it's sqrt(MAX)

thanx in advance....

January 8, 2010 12:48 AM



Zobayer Hasan said...

@Arun, I have added the description for the macros. I think you already got it as 1 and a half months passed and I was really busy. :(

March 10, 2010 12:44 AM

Anonymous said...

For anyone who is wondering what the macros are doing, shifting to right (c>>n) divides c by 2^n . So 32>>2 divides 32 by $2^2(=4)$, which results in 8. So, 32>>2=8. the bit shifting in those macros divides by 64 to make the address aligned with the word size of the machine(64-bits). The AND operator is a fast modulus operation for power of two numbers. As long as Y is some power of two(2^n , 2,4,8,16,32,64,128, etc...) X%Y can be written as X&Y-1. So, for example N%32 can be written as N&31. Wikipedia is a golden source of information, and google helps too.

September 13, 2010 9:48 PM

Anonymous said...

How do you print out the primes in the segment? Can you also post the main function of the program pls?

November 20, 2010 5:54 AM



Zobayer Hasan said...

@Anonymous, you see the last loop in segmented sieve function? That is counting number of primes in the interval [a, b] mapping a -> 0 and b -> b-a.

AdChriss

Log calls on Cisco UCME

Log via Radius or Syslog on UCME Free trial, advanced call reporting www.tri-line.com

Call Recording Systems

Fast & Easy To Use Call Recording Solutions. Get A Free Quote Online! www.element80.co....

<u>Software</u> <u>Developer Arnhem</u>

C++ Developer/ Java Developer Object oriented design www.hays.nl

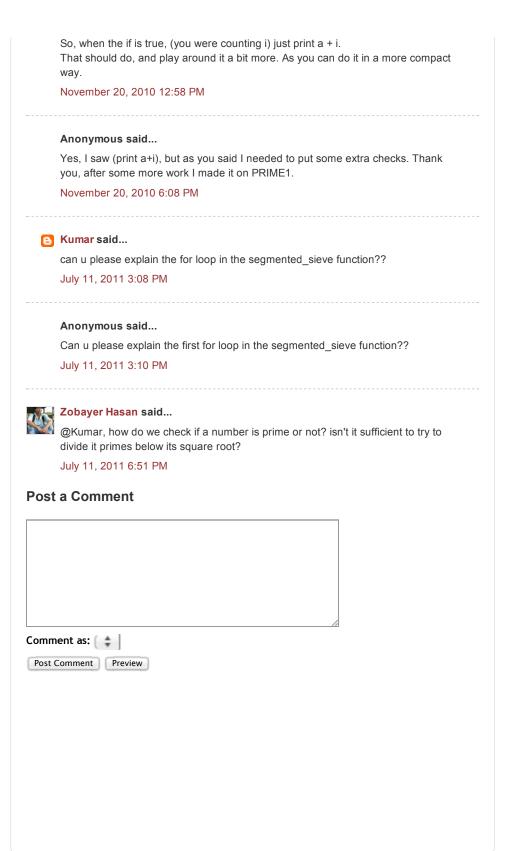
High Performance Servers

Delivered on PEER 1's SuperNetwork. 100% uptime SLA. Chat to us now! Peer1Hosting.co.uk...

1 Month Free IT Support

Limited Offer from Taribo Ltd To first 50 companies that apply www.taribo.co.uk







Newer Post Home Older Post

Subscribe to: Post Comments (Atom)

©2011 Zobayer Hasan. Picture Window template. Template images by Jason Morrow. Powered by Blogger.

I am only one, but still I am one.
I cannot do everything, but still I can do something.

And because I cannot do everything I will not refuse to do the something that I can do.

{Helen Keller}